

Bike Price Prediction using Linear Regression

Importing the libraries

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

Importing the dataset

```
df =
pd.read_csv('https://github.com/YBI-Foundation/Dataset/raw/main/Bike
%20Prices.csv')
```

Get Missing Values Drop

```
df=df.dropna()
```

Get Information of Dataframe

```
df.info() #gives column name, count, not null category, D-type(data
type)
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 626 entries, 0 to 625
Data columns (total 8 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Brand            626 non-null    object  
 1   Model             626 non-null    object  
 2   Selling_Price    626 non-null    int64   
 3   Year              626 non-null    int64   
 4   Seller_Type      626 non-null    object  
 5   Owner             626 non-null    object  
 6   KM_Driven        626 non-null    int64   
 7   Ex_Showroom_Price 626 non-null    float64 
dtypes: float64(1), int64(3), object(4)
memory usage: 44.0+ KB
```

```
df.describe() #gives the linear relation of each column with another
column
```

	Selling_Price	Year	KM_Driven	Ex_Showroom_Price
count	626.000000	626.000000	626.000000	6.260000e+02
mean	59445.164537	2014.800319	32671.576677	8.795871e+04
std	59904.350888	3.018885	45479.661039	7.749659e+04

```

min      6000.00000  2001.00000      380.00000      3.049000e+04
25%     30000.00000  2013.00000    13031.25000      5.485200e+04
50%     45000.00000  2015.00000    25000.00000      7.275250e+04
75%     65000.00000  2017.00000    40000.00000      8.703150e+04
max     760000.00000 2020.00000   585659.00000     1.278000e+06

```

```
df.head()
```

	Brand Owner \ owner	Model	Selling_Price	Year	Seller_Type	
0	TVS	TVS XL 100	30000	2017	Individual	1st
1	Bajaj	Bajaj ct 100	18000	2017	Individual	1st
2	Yo	Yo Style	20000	2011	Individual	1st
3	Bajaj	Bajaj Discover 100	25000	2010	Individual	1st
4	Bajaj	Bajaj Discover 100	24999	2012	Individual	2nd
	owner					

	KM_Driven	Ex_Showroom_Price
0	8000	30490.0
1	35000	32000.0
2	10000	37675.0
3	43000	42859.0
4	35000	42859.0

```
df.isnull().sum() #(df.isna().sum() gives same result)
#gives the sum of all null values columns-wise
```

```

Brand          0
Model          0
Selling_Price 0
Year           0
Seller_Type   0
Owner          0
KM_Driven     0
Ex_Showroom_Price 0
dtype: int64

```

```
df.nunique() #gives total no. of unique entries
```

Brand	18
Model	183
Selling_Price	99
Year	18
Seller_Type	2
Owner	4
KM_Driven	219

```
Ex_Showroom_Price      230
dtype: int64

df.shape

(626, 8)

#Get Categories and Counts of Categorical Variables
df[['Brand']].value_counts()

Brand
Honda        170
Bajaj         143
Hero          108
Yamaha        94
Royal          40
TVS            23
Suzuki         18
KTM             6
Mahindra        6
Kawasaki        4
UM              3
Activa          3
Harley           2
Vespa            2
BMW              1
Hyosung           1
Benelli           1
Yo               1
dtype: int64

df[['Model']].value_counts()

Model
Honda Activa [2000-2015]                23
Honda CB Hornet 160R                     22
Bajaj Pulsar 180                         20
Yamaha FZ S V 2.0                        16
Bajaj Discover 125                       16
                                         ..
Royal Enfield Thunderbird 500             1
Royal Enfield Continental GT [2013 - 2018] 1
Royal Enfield Classic Stealth Black       1
Royal Enfield Classic Squadron Blue        1
Yo Style                                     1
Length: 183, dtype: int64

df[['Seller_Type']].value_counts()

Seller_Type
Individual        623
```

```

Dealer      3
dtype: int64

df[['Owner']].value_counts()

Owner
1st owner    556
2nd owner     66
3rd owner      3
4th owner      1
dtype: int64

```

Get Column Names

```
df.columns
```

```

Index(['Brand', 'Model', 'Selling_Price', 'Year', 'Seller_Type',
'Owner',
       'KM_Driven', 'Ex_Showroom_Price'],
      dtype='object')

```

Get Encoding of Categorical Features

```

df.replace({'Seller_Type':{'Individual':0,'Dealer':1}},inplace=True)

df.replace({'Owner':{'1st owner':0,'2nd owner':1,'3rd owner':2,'4th
owner':3}},inplace=True)

#X=pd.get_dummies(X,columns=['Seller_Type','Owner'],drop_first=True)

```

```
df
```

	Brand	Model	Selling_Price	Year	\
0	TVS	TVS XL 100	30000	2017	
1	Bajaj	Bajaj ct 100	18000	2017	
2	Yo	Yo Style	20000	2011	
3	Bajaj	Bajaj Discover 100	25000	2010	
4	Bajaj	Bajaj Discover 100	24999	2012	
..
621	Harley	Harley-Davidson Street 750	330000	2014	
622	Kawasaki	Kawasaki Ninja 650 [2018-2019]	300000	2011	
623	Kawasaki	Kawasaki Ninja 650 [2018-2019]	425000	2017	
624	Suzuki	Suzuki GSX S750	760000	2019	
625	Harley	Harley-Davidson Street Bob	750000	2013	
	Seller_Type	Owner	KM_Driven	Ex_Showroom_Price	
0	0	0	8000	30490.0	
1	0	0	35000	32000.0	
2	0	0	10000	37675.0	
3	0	0	43000	42859.0	
4	0	1	35000	42859.0	

```

...
621      0      3     6500      534000.0
622      0      0    12000      589000.0
623      0      1    13600      599000.0
624      0      0     2800      752020.0
625      0      1    12000      1278000.0

```

[626 rows x 8 columns]

Define y

```
y=df['Selling_Price']
```

```
y.shape
```

(626,)

```
y
```

```

0      30000
1      18000
2      20000
3      25000
4      24999

```

```

...
621    330000
622    300000
623    425000
624    760000
625    750000

```

Name: Selling_Price, Length: 626, dtype: int64

```
X=df[['Year','Seller_Type','Owner','KM_Driven','Ex_Showroom_Price']]
```

```
X.shape
```

(626, 5)

```
X
```

	Year	Seller_Type	Owner	KM_Driven	Ex_Showroom_Price
0	2017	0	0	8000	30490.0
1	2017	0	0	35000	32000.0
2	2011	0	0	10000	37675.0
3	2010	0	0	43000	42859.0
4	2012	0	1	35000	42859.0
...
621	2014	0	3	6500	534000.0
622	2011	0	0	12000	589000.0
623	2017	0	1	13600	599000.0
624	2019	0	0	2800	752020.0
625	2013	0	1	12000	1278000.0

```
[626 rows x 5 columns]
```

Train Test Split

```
from sklearn.model_selection import train_test_split
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=
train_test_split(X,y,test_size=0.3,random_state=2323)
X_train.shape, X_test.shape, y_train.shape, y_test.shape
((438, 5), (188, 5), (438,), (188,))
```

Get Model Test

```
from sklearn.linear_model import LinearRegression
lr= LinearRegression()
lr.fit(X_train, y_train)
LinearRegression()
```

Get Model Prediction

```
y_pred= lr.predict(X_test)
y_pred.shape
(188,)

y_pred
array([ 35870.79961642,   65547.96126415,   12371.09961996,
       64148.8895839 ,
        7062.69388653,   54402.63506136,   -1912.44035159,
      29775.20161928,
        14243.08480753,   22080.83178243,   84025.43006156,
     38292.33063928,
        29354.69337304,   21525.66151869,   51718.07386067,
    133108.79218061,
        36717.94936617,   61355.37093641,   28387.67135344,
    43699.33828063,
        73829.7721122 ,  121281.37404139,  257025.1574888 , -
  13165.16881032,
        46855.75420218,   52525.81083583,   21528.77516102,
    58963.76859935,
        36453.23827565,   59092.23873301,   52537.87619985,
    31275.42528094,
```

50620.8218115 ,	67631.38902032,	56630.86444395,
117293.32518539,	34037.48760149,	49778.37056618,
21536.11351636,	111705.80584799,	
44867.13135661,	74909.17128765,	421105.83467059,
73670.58645457,		
51287.25308003,	14423.62915765,	99846.44490246,
54487.16853425,		
101946.69838917,	79576.38985325,	117362.89863505,
37565.98215828,		
68125.63747921,	35121.70517691,	52265.51803945,
54517.4438797 ,		
46366.66325235,	192872.63669555,	50035.40444538,
95226.6560928 ,		
50547.81489465,	25268.83224858,	62635.88716991,
34669.44918167,		
117369.04807865,	79570.16256859,	29637.7731142 ,
29989.47887016,		
7756.2225401 ,	29835.95697668,	67271.25283088,
44340.23851247,		
91684.46261394,	63810.8340198 ,	80449.25240847,
35102.22261493,		
43725.02582983,	61355.87846011,	32852.18571348,
47285.70027847,		
81095.64161965,	53340.80472393,	35507.12858801,
127940.00023025,		
54748.27265082,	33764.27590934,	21555.24112081,
47793.88009026,		
99600.47368742,	26427.56110399,	50843.0311352 ,
2393.36675538,		
44245.38463449,	30557.61972929,	64331.44738662,
46052.30174221,		
127948.19300162,	25875.99311672,	54013.85916036,
47862.33637501,		
54548.65103582,	45496.28368746,	31788.72526496,
25415.45959194,		
31554.54862708,	41029.82132688,	46789.06447165,
68798.88433921,		
118678.1466251 ,	38280.26527526,	51618.90192728,
60607.30748903,		
40486.80042414,	23634.458532 ,	48728.72961516,
133108.01377003,		
49844.92467094,	14273.22715688,	65610.31437702,
111888.54121412,		
71709.81448763,	140861.10089833,	29964.08104634,
49711.03130947,		
95879.34214946,	58311.43819235,	21723.74978983,
134163.14808106,		
132056.05882463,	55999.53310058,	40556.27936137,
71719.15541461,		

```
    61303.17569047, 69364.60182646, 51074.45481062,
28320.44222635,
    51773.58311705, 239835.53704415, 64318.1023156 ,
37399.68698658,
    38289.86152092, 59716.48021255, 21721.41455808,
14608.88730471,
    52467.53834686, 50611.39269332, 50026.06245928,
11788.30069698,
    59391.43797137, 51698.22439083, 59672.88921996,
38739.74209449,
    52091.16644551, 17201.62221796, 70016.38987149,
30835.15898424,
    101950.59044208, 46806.18950446, 60690.11143484,
126922.35408882,
    7089.9382569 , 84927.52216659, 109939.32315248,
60867.65085507,
    13311.86850318, 28477.46071707, 35870.79961642,
40632.64055195,
    130166.96427995, 68066.86002907, 34033.59554858,
54548.88455899,
    42884.61600956, 37735.67541855, 25863.9277527 ,
58185.3564429 ,
    138145.48035431, 112389.92337885, 60911.95193693,
50891.34804191,
    121286.7450744 , 30918.49434596, 51936.08269254,
12699.70125396,
    57749.07933698, 45996.1445545 , 23637.15540436,
35095.67384671])
```

Get Model Evaluation

```
from sklearn.metrics import mean_squared_error, mean_absolute_error,
mean_absolute_percentage_error, r2_score

mean_squared_error(y_test,y_pred)
293014701.66187763

mean_absolute_percentage_error(y_test,y_pred)
0.24947046843526693

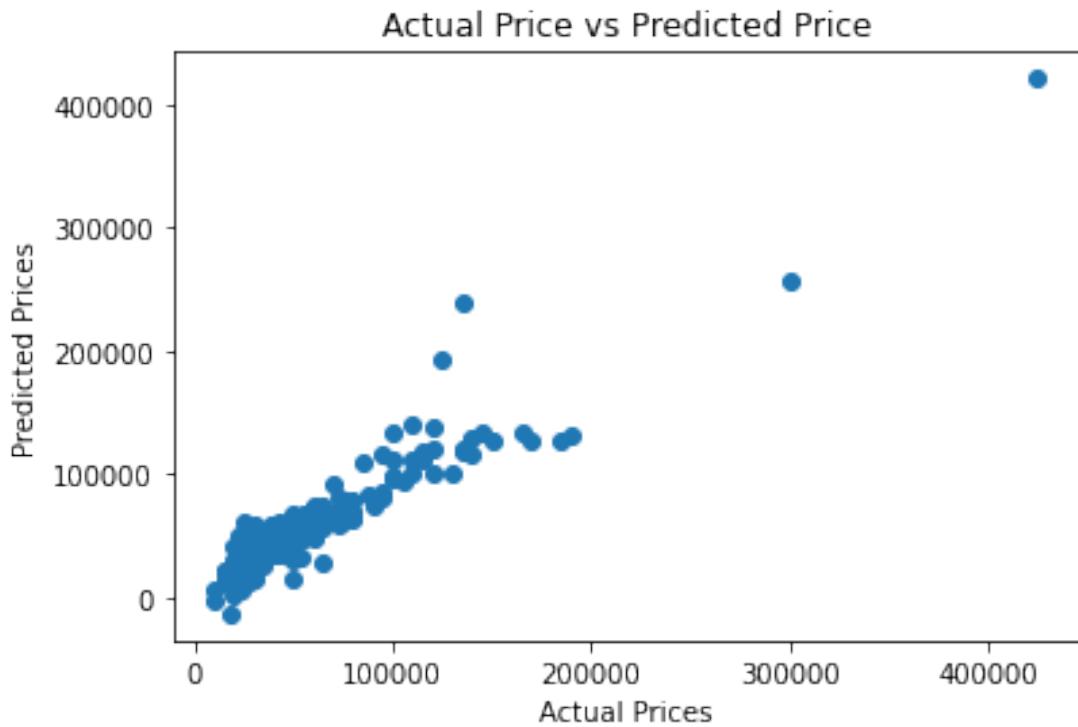
r2_score(y_test,y_pred)
0.8738616789555578

#Get Visualization of Actual Vs Predicted Results

import matplotlib.pyplot as plt

plt.scatter(y_test,y_pred)
plt.xlabel("Actual Prices")
```

```
plt.ylabel('Predicted Prices')
plt.title('Actual Price vs Predicted Price')
plt.show()
```



```
X_new= df.sample(1)
```

```
X_new
```

```
    Brand           Model  Selling_Price  Year Seller_Type  Owner
4  Bajaj   Bajaj Discover 100          24999  2012            0      1
```

```
    KM_Driven  Ex_Showroom_Price
4      35000          42859.0
```

```
X_new.shape
```

```
(1, 8)
```

```
X_new=X_new.drop(['Selling_Price', 'Brand', 'Model'],axis=1)
```

```
X_new
```

```
    Year Seller_Type  Owner  KM_Driven  Ex_Showroom_Price
4  2012          0      1      35000          42859.0
```

```
X_new.shape
```

(1, 5)

```
y_pred_new= lr.predict(X_new)
y_pred_new
array([11014.37068462])
```