

Importing the libraries

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

Importing the dataset

```
df =
pd.read_csv('https://github.com/YBI-Foundation/Dataset/raw/main/MPG.csv')
```

Get Information of Dataframe

```
df.info() #gives column name, count, not null category, D-type(data type)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   mpg              398 non-null    float64
 1   cylinders        398 non-null    int64  
 2   displacement     398 non-null    float64
 3   horsepower       392 non-null    float64
 4   weight            398 non-null    int64  
 5   acceleration     398 non-null    float64
 6   model_year       398 non-null    int64  
 7   origin            398 non-null    object 
 8   name              398 non-null    object 
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB
```

```
df.describe() #gives the linear relation of each column with another column
```

```
               #if the mean of the columns not of same power ,we need to standardize data
```

	mpg	cylinders	displacement	horsepower	
weight \ count	398.000000	398.000000	398.000000	392.000000	398.000000
mean	23.514573	5.454774	193.425879	104.469388	2970.424623
std	7.815984	1.701004	104.269838	38.491160	846.841774

min	9.000000	3.000000	68.000000	46.000000	1613.000000
25%	17.500000	4.000000	104.250000	75.000000	2223.750000
50%	23.000000	4.000000	148.500000	93.500000	2803.500000
75%	29.000000	8.000000	262.000000	126.000000	3608.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000

	acceleration	model_year
count	398.000000	398.000000
mean	15.568090	76.010050
std	2.757689	3.697627
min	8.000000	70.000000
25%	13.825000	73.000000
50%	15.500000	76.000000
75%	17.175000	79.000000
max	24.800000	82.000000

```
df.head(13)
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	18.0	8	307.0	130.0	3504	12.0	
1	15.0	8	350.0	165.0	3693	11.5	
2	18.0	8	318.0	150.0	3436	11.0	
3	16.0	8	304.0	150.0	3433	12.0	
4	17.0	8	302.0	140.0	3449	10.5	
5	15.0	8	429.0	198.0	4341	10.0	
6	14.0	8	454.0	220.0	4354	9.0	
7	14.0	8	440.0	215.0	4312	8.5	
8	14.0	8	455.0	225.0	4425	10.0	
9	15.0	8	390.0	190.0	3850	8.5	
10	15.0	8	383.0	170.0	3563	10.0	
11	14.0	8	340.0	160.0	3609	8.0	
12	15.0	8	400.0	150.0	3761	9.5	

	model_year	origin	name		
0	70	usa	chevrolet	chevelle	malibu
1	70	usa	buick	skylark	320
2	70	usa	plymouth	satellite	
3	70	usa	amc	rebel	sst
4	70	usa	ford	torino	
5	70	usa	ford	galaxie	500
6	70	usa	chevrolet	impala	
7	70	usa	plymouth	fury	iii
8	70	usa	pontiac	catalina	
9	70	usa	amc	ambassador	dpl
10	70	usa	dodge	challenger	se

```

11      70    usa      plymouth 'cuda 340
12      70    usa      chevrolet monte carlo

df.isnull().sum() #(df.isna().sum() gives same result)
#gives the sum of all null values columns-wise

mpg          0
cylinders    0
displacement 0
horsepower   6
weight        0
acceleration 0
model_year   0
origin        0
name          0
dtype: int64

df.nunique() #gives total no. of unique entries

mpg          129
cylinders    5
displacement 82
horsepower   93
weight        351
acceleration 95
model_year   13
origin        3
name          305
dtype: int64

df.columns #give column names in the dataframe
Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
       'acceleration', 'model_year', 'origin', 'name'],
      dtype='object')

df.shape
(392, 9)

df.corr()      #as mpg cylinders and displacement have very high
correlation we will use only one of those
#in this case we are using displacement

\           mpg  cylinders  displacement  horsepower  weight
mpg      1.000000 -0.777618   -0.805127 -0.778427 -0.832244
cylinders -0.777618  1.000000    0.950823  0.842983  0.897527
displacement -0.805127   0.950823     1.000000  0.897257  0.932994

```

```
horsepower -0.778427  0.842983  0.897257  1.000000  0.864538
weight      -0.832244  0.897527  0.932994  0.864538  1.000000
acceleration 0.423329 -0.504683 -0.543800 -0.689196 -0.416839
model_year   0.580541 -0.345647 -0.369855 -0.416361 -0.309120
```

```
          acceleration model_year
mpg           0.423329  0.580541
cylinders     -0.504683 -0.345647
displacement  -0.543800 -0.369855
horsepower    -0.689196 -0.416361
weight        -0.416839 -0.309120
acceleration  1.000000  0.290316
model_year    0.290316  1.000000
```

```
#Remove Missing Values
```

```
df=df.dropna()
```


```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 392 entries, 0 to 397
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   mpg         392 non-null    float64
 1   cylinders   392 non-null    int64  
 2   displacement 392 non-null    float64
 3   horsepower   392 non-null    float64
 4   weight       392 non-null    int64  
 5   acceleration 392 non-null    float64
 6   model_year   392 non-null    int64  
 7   origin       392 non-null    object 
 8   name         392 non-null    object 
dtypes: float64(4), int64(3), object(2)
memory usage: 30.6+ KB
```


```
#Data Visualization
```

```
sns.pairplot(df,x_vars=['displacement','horsepower','weight','acceleration'],y_vars='mpg')
```

```
<seaborn.axisgrid.PairGrid at 0x7fef3f9fc850>
```



```
sns.regplot(x='displacement', y='mpg', data=df)      #regression line
<matplotlib.axes._subplots.AxesSubplot at 0x7fef3d0464d0>
```



```
#Define Target Variable y and Feature X
df.columns
Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
       'acceleration', 'model_year', 'origin', 'name'],
      dtype='object')
y=df['mpg']
y.shape
(392,)
X=df[['displacement', 'horsepower', 'weight', 'acceleration']]
X.shape
```

(392, 4)

```
#Scaling Data
from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
X= ss.fit_transform(X)

X
array([[ 1.07728956,   0.66413273,   0.62054034,  -1.285258  ],
       [ 1.48873169,   1.57459447,   0.84333403,  -1.46672362],
       [ 1.1825422 ,   1.18439658,   0.54038176,  -1.64818924],
       ...,
       [-0.56847897,  -0.53247413,  -0.80463202,  -1.4304305 ],
       [-0.7120053 ,  -0.66254009,  -0.41562716,   1.11008813],
       [-0.72157372,  -0.58450051,  -0.30364091,   1.40043312]])]

pd.DataFrame(X).describe()

          0            1            2            3
count  3.920000e+02  3.920000e+02  3.920000e+02  3.920000e+02
mean   -2.537653e-16 -4.392745e-16  5.607759e-17  6.117555e-16
std    1.001278e+00  1.001278e+00  1.001278e+00  1.001278e+00
min    -1.209563e+00 -1.520975e+00 -1.608575e+00 -2.736983e+00
25%   -8.555316e-01 -7.665929e-01 -8.868535e-01 -6.410551e-01
50%   -4.153842e-01 -2.853488e-01 -2.052109e-01 -1.499869e-02
75%   7.782764e-01  5.600800e-01  7.510927e-01  5.384714e-01
max    2.493416e+00  3.265452e+00  2.549061e+00  3.360262e+00
```

Splitting the dataset into the Training set and Test set

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.3, random_state = 1)

X_train.shape, X_test.shape, y_train.shape, y_test.shape
((274, 4), (118, 4), (274,), (118,))
```

LINEAR REGRESSION

```
from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(X_train,y_train)

LinearRegression()
```

```
lr.intercept_
23.19931701979532

lr.coef_
array([-0.71568189, -1.1167925 , -4.51963105,  0.41172277])

mpg= 23.2 +(-0.71displacement) + (-1.11horsepower) + (-4.51weight) + (-0.41acceleration)
```

Model Prediction

```
y_pred=lr.predict(X_test)

y_pred.shape
(118,)

y_pred

array([26.28276367, 29.08833366, 28.17888013, 24.06092615,
       30.32940196,
       29.34559658, 29.47244419, 14.16480917, 28.12054088,
       28.62243268,
       21.19699253, 29.00202854, 17.08446147, 30.6393795 ,
       26.68062515,
       18.16545075, 27.05814569, 30.19749225,  9.75469894,
       23.47905709,
       28.25978712, 19.48402464, 18.1064714 , 13.28502247,
       10.78371487,
       15.23539746, 30.32970958, 19.85446483, 22.37999549,
       26.96328669,
       18.53096933, 25.30938579, 12.23913664, 23.35707834,
       20.75265097,
       14.43741584, 19.4800578 , 20.51028037, 30.30992209,
       28.70496355,
       11.62869958, 11.71491907, 23.55687348, 23.54831788,
       23.99559915,
       19.11133191,  9.47947187, 30.64274232, 21.73778331,
       8.26480713,
       16.82257963, 22.88819913, 25.65631731, 27.24852928,
       31.15303313,
       23.38685254, 24.44456186, 25.87613663, 25.81319389,
       31.60155856,
       23.32071279, 27.06454436, 31.01331394, 20.51927114,
       23.09760801,
       22.33665492, 23.57958269, 16.09819394, 30.27109569,
       8.5745448 ,
       26.89318004, 18.68219122, 16.65807253, 28.88072244,
       26.96138898,
```

```
    30.06105048, 16.507612 , 12.02031005, 12.93921033,
26.82226838,
    19.87853884, 28.6680564 , 30.32376371, 28.69240544,
14.84426249,
    20.49630205, 30.74330121, 25.97075053, 31.06972194,
23.9671303 ,
    28.82409624, 21.11682718, 14.60616423, 11.49977764,
7.75259844,
    24.12423109, 14.48534835, 24.39258351, 23.40743093,
13.20487542,
    28.58355011, 26.42326644, 23.31389675, 28.91626441,
28.95276995,
    29.56630042, 23.43332526, 31.38563943, 14.14444937,
18.34692392,
    30.74367965, 23.9993707 , 30.32374274, 27.61841915,
15.3921527 ,
    25.62422581, 20.97391756, 12.8109788 ])
```

Model Evaluation

```
from sklearn.metrics import mean_squared_error, mean_absolute_error,
mean_absolute_percentage_error, r2_score

mean_absolute_error(y_test,y_pred)
3.1259768036439923

mean_absolute_percentage_error(y_test,y_pred)
0.13268990973331238

r2_score(y_test,y_pred)           #regressor.score(x_test,y_test)    ---
gives same result
0.7209702318140772

#Polynomial Regression

from sklearn.preprocessing import PolynomialFeatures
poly=PolynomialFeatures(degree=2,
interaction_only=True,include_bias=False)

X_train2=poly.fit_transform(X_train)
X_test2=poly.fit_transform(X_test)
lr.fit(X_train2,y_train)
LinearRegression()
lr.intercept_
20.98947080512586
```

```
lr.coef_
array([-2.90070836, -4.57930091, -1.18400444, -0.48222393,
2.10945694,
       0.59137383, -0.46726197, -0.32012722, -0.02876052,
0.56719103])

y_pred_poly=lr.predict(X_test2)
mean_absolute_error(y_test,y_pred_poly)
2.8067976869672058
mean_absolute_percentage_error(y_test,y_pred_poly)
0.11485347284477058
r2_score(y_test,y_pred_poly)
0.7657248979487836
```